

## Experiment 4: Map-Reduce Program

### **Program:-**

```
package supreetha;

import java.io.IOException;
import java.util.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
public class TransactionCount {
    //MAPPER CODE
    public static class Map extends MapReduceBase implements
    Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        //private Text word = new Text();
        public void map(LongWritable key, Text value, OutputCollector<Text,
        IntWritable> output, Reporter reporter) throws IOException {
            String myString = value.toString();
            String[] userCount = myString.split(",");
            output.collect(new Text(userCount[3]), one);
        }
    }
    //REDUCER CODE
    public static class Reduce extends MapReduceBase implements
    Reducer<Text, IntWritable, Text, IntWritable> {
        public void reduce(Text key, Iterator<IntWritable> values,
                           OutputCollector<Text, IntWritable> output, Reporter reporter) throws
                           IOException { //little: {1,1}
            int finaluserCount = 0 ;
            Text mykey = key ;
            while(values.hasNext()) {
                IntWritable value = values.next();
                finaluserCount += value.get();
            }
            output.collect(mykey, new IntWritable(finaluserCount));
        }
    }
    //DRIVER CODE
    public static void main(String[] args) throws Exception {
        JobConf conf = new JobConf(TransactionCount.class);
        conf.setJobName("transactioncount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(Map.class);
        conf.setCombinerClass(Reduce.class);
    }
}
```

```

        conf.setReducerClass(Reduce.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        JobClient.runJob(conf);
    }
}

```

```

1 package supreetha;
2
3 import java.io.IOException;
4 import java.util.*;
5 import org.apache.hadoop.fs.Path;
6 import org.apache.hadoop.io.*;
7 import org.apache.hadoop.mapred.*;
8 public class TransactionCount {
9     //MAPPER CODE
10    public static class Map extends MapReduceBase implements
11        Mapper<LongWritable, Text, Text, IntWritable> {
12        private final static IntWritable one = new IntWritable(1);
13        //private Text word = new Text();
14        public void map(LongWritable key, Text value, OutputCollector<Text,
15        IntWritable> output, Reporter reporter) throws IOException {
16            String myString = value.toString();
17            String[] userCount = myString.split(",");
18            output.collect(new Text(userCount[3]), one);
19        }
20    }
21    //REDUCER CODE
22    public static class Reduce extends MapReduceBase implements
23        Reducer<Text, IntWritable, Text, IntWritable> {

```

### .csv File:

	A	B	C	D	E	F
1	sl no	card name	user name	amount		
2		1 abc	mno	1000		
3		2 def	pqr	1000		
4		3 ghi	stu	2000		
5		4 jkl	vwx	2000		
6						
7						
8						
9						
10						

### Commands:

\$ cd \$HADOOP\_HOME

```
$ cd sbin  
./start-all.sh
```

```
hadoop@admin1-HP-280-G4-MT-Business-PC:~$ cd $HADOOP_HOME
```

```
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1$ cd sbin  
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ start-all.sh  
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.  
WARNING: This is not a recommended production deployment configuration.  
WARNING: Use CTRL-C to abort.  
Starting namenodes on [localhost]
```

```
hdfs dfs -mkdir -p /supreetha  
hdfs dfs -copyFromLocal /home/hadoop/Desktop/1nt19is166_supreetha/exp_4.csv  
/supreetha
```

```
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -mkdir -p /supreetha  
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -copyFromLocal /home/hadoop/Desktop/1nt19is166_supreetha/exp_4.csv /supreetha  
2022-06-23 14:09:46,882 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
```

```
hadoop jar /home/hadoop/Desktop/supreetha1.jar /supreetha /supreethaop
```

```
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hadoop jar /home/hadoop/Desktop/supreetha1.jar /supreetha /supreethaop  
2022-06-23 14:11:05,288 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032  
2022-06-23 14:11:05,415 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032  
2022-06-23 14:11:05,549 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.  
2022-06-23 14:11:05,573 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1655973226068_0001  
2022-06-23 14:11:05,680 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false  
2022-06-23 14:11:05,780 INFO mapred.FileInputFormat: Total input files to process : 1  
2022-06-23 14:11:05,830 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false  
2022-06-23 14:11:05,873 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false  
2022-06-23 14:11:05,897 INFO mapreduce.JobSubmitter: number of splits:2  
2022-06-23 14:11:06,039 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false  
2022-06-23 14:11:06,055 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1655973226068_0001  
2022-06-23 14:11:06,055 INFO mapreduce.JobSubmitter: Executing with tokens: []
```

```
2022-06-21 09:31:18,662 INFO mapreduce.Job: Job job_1655783592787_0001 running in uber mode : false  
2022-06-21 09:31:18,664 INFO mapreduce.Job: map 0% reduce 0%  
2022-06-21 09:31:22,742 INFO mapreduce.Job: map 100% reduce 0%  
2022-06-21 09:31:25,763 INFO mapreduce.Job: map 100% reduce 100%  
2022-06-21 09:31:26,782 INFO mapreduce.Job: Job job_1655783592787_0001 completed successfully  
2022-06-21 09:31:26,838 INFO mapreduce.Job: Counters: 54
```

```
hdfs dfs -cat /supreethaop/part*
```

```
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -cat /supreethaop/part*  
2022-06-23 14:12:21,902 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false  
1000 2  
2000 2  
amount 1  
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$
```